

# Kubernetes Networking



Owned by [Raghavendra Addanki](#) ...  
Last updated: May 15, 2023

## Kubernetes Internal Cluster Networking

- Kubernetes creates an Internal Cluster with a specific IP Subnet.
- The IPs in the cluster are accessible by all PODs in the Cluster.

### How to find this Cluster IP Range

```
kubectl describe cm kubeadm-config -n kube-system |grep Subnet
```

```
1 linuxadmin@master:~$ kubectl describe cm kubeadm-config -n kube-system |grep Subnet
2   serviceSubnet: 10.96.0.0/12
```

- When you create a Kubernetes Cluster, The Master Node that runs the API Server, will take the First available IP and sets that IP as Kubernetes Cluster IP.
- If any Node or POD want to communicate with Kubernetes API Server, they will use this IP.

### How to see the Kubernetes Default Cluster IP

```
kubectl get service --all-namespaces
```

```
1 linuxadmin@master:~$ kubectl get service --all-namespaces
2 NAMESPACE      NAME           TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
3 default         kubernetes     ClusterIP     10.96.0.1     <none>         443/TCP          179m
4 kube-system     kube-dns       ClusterIP     10.96.0.10    <none>         53/UDP, 53/TCP, 9153/TCP 179m
5
```

- As you can see kubernetes API Service is using IP Address: `ClusterIP 10.96.0.1`

### Internally this IP is mapped to the actual Master Node IP:

```
kubectl get endpoints kubernetes
```

```
1 linuxadmin@master:~$ kubectl get endpoints kubernetes
2 NAME           ENDPOINTS          AGE
3 kubernetes     192.168.1.50:6443  3h7m
```

## Kubernetes POD Networking

- POD have their own IP Subnet Ranges.
- They are not part of this Kubernetes Service Subnet: `serviceSubnet: 10.96.0.0/12`
- Each Worker Node is allotted a POD Network Range.
- When a POD is created on a particular Worker Node, The POD inherits one of the IP in the POD IP Range.

### How to find POD IP Range

```
kubectl get ipamblocks.crd.projectcalico.org
```

```
1 linuxadmin@master:~$ kubectl get ipamblocks.crd.projectcalico.org
2 NAME           AGE
3 172-16-104-0-26  57m
4 172-16-166-128-26  57m
5 172-16-219-64-26   58m
```

- But we the above output did not tell us which POD has got what IP Range.

### How to find individual Worker Node POD IP Range

```
kubectl get ipamblocks.crd.projectcalico.org -o jsonpath="{range .items[*]}{'podNetwork: '}{.spec.cidr}{'\t NodeIP: '}{.spec.affinity}{'\n'}"
```

```
1 linuxadmin@master:~$ kubectl get ipamblocks.crd.projectcalico.org -o jsonpath="{range .items[*]}{'podNetwork: '}{
2 podNetwork: 172.16.104.0/26      NodeIP: host:node2
3 podNetwork: 172.16.166.128/26    NodeIP: host:node1
4 podNetwork: 172.16.219.64/26     NodeIP: host:master
```

- Any POD created on **node1** will get an IP Address in the Range: 172.16.166.128/26

```
podNetwork: 172.16.166.128/26      NodeIP: host:node1
```

- Any POD created on **node2** will get an IP Address in the Range: 172.16.104.0/26

```
podNetwork: 172.16.104.0/26      NodeIP: host:node2
```

---

### Create some PODs and verify

#### First POD:

```
sudo nano ssl-website1.yaml
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-ssl-website1
5 spec:
6   containers:
7     - name: con-ssl-website1
8       image: tanvisinghny/ssl-website
9       ports:
10        - containerPort: 80
11        - containerPort: 443
```

```
1 kubectl create -f ssl-website1.yaml
```

#### Second POD:

```
sudo nano ssl-website2.yaml
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-ssl-website2
5 spec:
6   containers:
7     - name: con-ssl-website2
8       image: tanvisinghny/ssl-website
9       ports:
10        - containerPort: 80
11        - containerPort: 443
```

```
1 kubectl create -f ssl-website2.yaml
```

---

## Get the Individual POD IPs

```
1 linuxadmin@master:~$ kubectl get pods -o wide
2 NAME                READY   STATUS    RESTARTS   AGE   IP              NODE   NOMINATED NODE   READINESS GATES
3 pod-ssl-website1    1/1     Running   0           16s   172.16.166.130  node1   <none>           <none>
4 pod-ssl-website2    1/1     Running   0           11s   172.16.104.2   node2   <none>           <none>
```

POD on node1 gets IP address: 172.16.166.130

POD on node2 gets IP address: 172.16.104.2

Compare with POD IP Range we discussed earlier:

```
1 podNetwork: 172.16.104.0/26      NodeIP: host:node2
2 podNetwork: 172.16.166.128/26    NodeIP: host:node1
```

---

## Go Inside the POD

```
1 kubectl exec --stdin --tty pod-ssl-website1 -- /bin/bash
```

Kubernetes will inject the Kubernetes Cluster IP and Port Number into the Environment Variables of every POD, so that Each POD can communicate with Kubernetes API Service:

Within the POD run the `printenv` command:

```
1 root@pod-ssl-website1:/# printenv
2
3 ---
4 KUBERNETES_PORT_443_TCP_PROTO=tcp
5 KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1
6 KUBERNETES_SERVICE_HOST=10.96.0.1
7 KUBERNETES_PORT=tcp://10.96.0.1:443
8 KUBERNETES_PORT_443_TCP_PORT=443
```

That is how a POD knows How to talk to Kubernetes API Server.

---

## Exposing the Container to External World

- By default PODs and the containers inside are accessible only within POD Network.

While still inside the **POD1** we connected earlier try to access the Website hosted on **POD2**: 172.16.104.2

```
1 apt update
2 apt install curl -y
3
4 curl http://172.16.104.2
```

**POD-to-POD Communication successful!**

## But How to access this Website from External World?

- By creating a Service Object

Creating a Service Object that acts as a Broker between POD-network > Cluster-network > External-network

```
1 kubectl expose pod pod-ssl-website1 --type=NodePort --port=443 --target-port=443
```

## Error

```
1 linuxadmin@master:~$ kubectl expose pod pod-ssl-website1 --type=NodePort --port=443 --target-port=443
2 error: couldn't retrieve selectors via --selector flag or introspection: the pod has no labels and cannot be exposed
```

Without Label Service Object can not be created.

### Find Labels of both the PODS:

```
1 kubectl get pod pod-ssl-website1 -n default --show-labels
2
3 linuxadmin@master:~$ kubectl get pod pod-ssl-website1 -n default --show-labels
4 NAME                READY   STATUS    RESTARTS   AGE   LABELS
5 pod-ssl-website1    1/1     Running   0           29m   <none>
6
7 linuxadmin@master:~$ kubectl get pod pod-ssl-website2 -n default --show-labels
8 NAME                READY   STATUS    RESTARTS   AGE   LABELS
9 pod-ssl-website2    1/1     Running   0           29m   <none>
```

NOTE: We will discuss LABELS in next topic, but assigning a LABEL to a POD is mandatory.

### Label both the PODS:

```
1 linuxadmin@master:~$ kubectl label pods pod-ssl-website2 -n default name=app2
2 pod/pod-ssl-website2 labeled
3
4 linuxadmin@master:~$ kubectl label pods pod-ssl-website1 -n default name=app1
5 pod/pod-ssl-website1 labeled
```

### Show the labels:

```
1 kubectl get pods --show-labels
2
3 linuxadmin@master:~$ kubectl get pods --show-labels
4 NAME                READY   STATUS    RESTARTS   AGE   LABELS
5 pod-ssl-website1    1/1     Running   0           36m   name=app1
6 pod-ssl-website2    1/1     Running   0           36m   name=app2
```

### Now create a Kubernetes Service that will expose the PODS to External World:

Exposing the First POD only to test.

```
1 kubectl expose pod pod-ssl-website1 --type=NodePort --name=svc-pod-ssl-website1
2
3 linuxadmin@master:~$ kubectl expose pod pod-ssl-website1 --type=NodePort --name=svc-pod-ssl-website1
4 service/svc-pod-ssl-website1 exposed
```

### Find the newly created Service:

```
1 linuxadmin@master:~$ kubectl get services
2 NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)                AGE
3 kubernetes          ClusterIP   10.96.0.1     <none>         443/TCP                4h6m
4 svc-pod-ssl-website1 NodePort     10.105.83.26  <none>         80:30686/TCP,443:30727/TCP 7s
```

```
svc-pod-ssl-website1  NodePort     10.105.83.26  <none>         80:30686/TCP,443:30727/TCP  7s
```

That is not Enough!

You can see that External IP is <none>.

### Patch the Service with the node1 physical IP = 192.168.1.51

```
kubectl patch svc svc-pod-ssl-website1 -n default -p '{"spec": {"type": "NodePort", "externalIPs":  
["192.168.1.51"]}}'
```

```
1 linuxadmin@master:~$ kubectl patch svc svc-pod-ssl-website1 -n default -p '{"spec": {"type": "NodePort", "externalIPs":  
2 service/svc-pod-ssl-website1 patched
```

**Again find Service details:**

```
1 kubectl get services  
2  
3 linuxadmin@master:~$ kubectl get services  
4 NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)                AGE  
5 kubernetes           ClusterIP    10.96.0.1      <none>          443/TCP                 4h11m  
6 svc-pod-ssl-website1 NodePort     10.105.83.26   192.168.1.51   80:30686/TCP,443:30727/TCP 5m7s  
7
```

Now the External IP is reflected.

---

**Access the Website now!**